

Assignment 2 — Selection Sort (Student B)

Technical Report (for submission & defense)

1) Project Overview

This project delivers a clean, production-ready **Java/Maven** implementation of **Selection Sort** with:

- **Two variants:** classic *standard* and *double-ended* (min & max per pass).
- **Early-exit optimization:** stops when the array is already sorted.
- **Metrics collection:** comparisons, swaps, reads, writes, iterations, early-termination flag.
- **CLI benchmark runner** and **JUnit 5 tests**.

The goal is to demonstrate algorithmic correctness, measure basic performance characteristics, and provide a reproducible benchmark harness.

2) Objectives & Requirements

- Implement Selection Sort (Student B role) with neat code and zero IDE errors.
 - Add early-exit optimization and (optionally) a double-ended pass.
 - Provide a small metrics utility to quantify operations.
 - Supply unit tests and a runnable CLI for demos/experiments.
-

3) Repository Structure

```
src/  
  main/  
    java/  
      algorithms/SelectionSort.java  
      metrics/PerformanceTracker.java  
      cli/BenchmarkRunner.java  
  test/  
    java/  
      algorithms/SelectionSortTest.java  
pom.xml
```

4) Build, Test, Run

- **Build & Tests**
- `mvn -q -DskipTests=false test`
- **Run CLI (single default run if no args)**
- # Example manual run with arguments:
- `mvn -q exec:java -Dexec.args="20 random standard true"`
- # Format: <size> <mode> <algo> <earlyExit> [seed]
- # mode: random|sorted|reversed|duplicates
- # algo: standard|double

- **IDE:** Run `cli.BenchmarkRunner.main()`; without args it performs one default random run.

5) Algorithm Design

5.1 Standard Selection Sort

- Outer loop index i scans from 0 to $n-2$.
- Find the index of the **minimum** element on $[i..n-1]$.
- Swap it with $a[i]$.
- **Loop invariant:** after the i -th pass, positions $0..i$ hold the $i+1$ smallest elements in sorted order.

5.2 Double-Ended Variant

- Maintain two pointers: `left` and `right`.
- In **one pass**, scan $[left..right]$ to find both **min** and **max**.
- Place `min` \rightarrow `left`, `max` \rightarrow `right`, then shrink to $left+1..right-1$.
- This reduces the **number of outer passes** roughly by half, while overall time is still $\Theta(n^2)$.

5.3 Early-Exit Optimization

- After each pass, check if the array is already non-decreasing.
If yes \rightarrow **terminate** (best-case $\Theta(n)$).

6) Complexity & Trade-offs

Variant	Best Case	Average / Worst Space
Standard + Early Exit	$\Theta(n)$ (already sorted)	$\Theta(n^2)$ $O(1)$
Double-Ended + Early	$\Theta(n)$ (already sorted)	$\Theta(n^2)$ $O(1)$

Notes

- Selection Sort is **not stable** (swaps can reorder equals). Stability would require extra logic (e.g., stable selection via shifting), which increases writes.
- Double-ended variant reduces *passes* but remains $O(n^2)$ because each pass still scans a linear range.

7) Correctness Argument (sketch)

- **Initialization:** Before any pass, the left prefix is empty; trivially sorted and minimal.
- **Maintenance:** Each pass selects the global minimum from the unsorted suffix and places it at the boundary (and, in the double-ended case, also places the global maximum at the

right boundary). The invariant is preserved: the prefix (and suffix) are sorted and contain the correct extremal elements.

- **Termination:** After $n-1$ (standard) or $\approx n/2$ (double-ended) passes, the unsorted range is empty; the array is fully sorted.

8) Implementation Details

8.1 `algorithms.SelectionSort`

- Public API:
- `SelectionSort.sort(int[] a,`
• `PerformanceTracker trackerOrNull,`
• `boolean earlyExit,`
• `boolean doubleEnded);`
- Input validation: `null` \rightarrow `IllegalArgumentException`; arrays of length < 2 return immediately.
- **Metrics hooks** around comparisons, swaps, and array accesses.
- `isSorted(..)` used for early exit.

8.2 `metrics.PerformanceTracker`

- Fields: `comparisons`, `swaps`, `arrayReads`, `arrayWrites`, `iterations`, `earlyTerminated`.
- Methods: `increment/add` counters, `toString()`, `csvHeader()`, `toCsvRow()`.

8.3 `cli.BenchmarkRunner`

- Parses arguments and runs one benchmark; if no args, performs a default single random run.
- Prints input (truncated), sorted output (truncated), human-readable metrics, and a one-line CSV.

8.4 Tests (`SelectionSortTest`)

- Edge cases: empty and single-element arrays.
- Duplicates handling.
- Random arrays compared against `Arrays.sort`.
- Early exit triggers on already sorted input.

9) Metrics & Experimental Notes

What is measured

- `comparisons`: every `a[j]` vs `a[minIdx]` (and `a[maxIdx]` in double-ended).
- `swaps`: swap operations.
- `arrayReads/arrayWrites`: rough counts for memory touches (instrumented estimates).
- `iterations`: outer passes (or left/right rounds).
- `earlyTerminated`: boolean indicating early exit.

Example CLI runs (typical)

```
=== Run ===
mode=random, algo=standard, earlyExit=true, n=20, seed=...
Input:  [ .... ]
Sorted: [ .... ]
Metrics: PerformanceTracker{comparisons=190, swaps=~19, arrayReads=...,
arrayWrites=..., iterations=19, earlyTerminated=false}
CSV:
comparisons,swaps,arrayReads,arrayWrites,iterations,earlyTerminated
190,19,.....,.....,19,false
```

Values depend on input and variant; early-exit becomes `true` for already sorted input, significantly reducing comparisons and iterations.

Interpretation

- On random data, Selection Sort performs $\sim n(n-1)/2$ comparisons ($\Theta(n^2)$) regardless of early exit.
 - On sorted input, early exit reduces work to $\Theta(n)$, which you can show by running with `mode=sorted`.
-

10) Edge Cases & Robustness

- **Empty / single-element arrays:** no action, metrics remain near zero.
 - **Large duplicates:** both variants handle equal keys; result is correct but not stable.
 - **Integer range:** uses `int`; for other primitives or generics, the design can be templated.
-

11) Limitations

- Not stable in current form.
 - $\Theta(n^2)$ time makes it unsuitable for very large datasets (demonstration/teaching use).
 - Early exit adds a linear check per pass (negligible vs overall $\Theta(n^2)$, but worth noting).
-

12) Possible Extensions

- **Stable selection sort** (shift instead of swap).
 - **Generic comparator-based API** (`Comparator<T>` with `T[]`).
 - **Micro-benchmarks** with JMH (Java Microbenchmark Harness).
 - Export metrics to CSV files or plot charts.
-

13) Conclusion

The project fulfills the assignment goals: correct Selection Sort implementations (standard & double-ended), early-exit optimization, measurable metrics, unit tests, and a straightforward CLI

benchmark. The code is modular, easy to run, and suitable for in-class demos and basic empirical analysis.